

/01




Project on Facebook Comment Volume Dataset

Python for data analysis

Martin Liger, Jean Baptiste Martin and Lamiae Maliki

Data overview



Data cleaning



**Creation of the
model**



Bonus work and Conclusion



Summary

/02

/03

Goal : predict how many comments a post will receive
overview of the dirty dataset:

```
[4] path_file = "/content/sample_data/Features_Variant_1.csv"  
data = pd.read_csv(path_file)  
print(data.head(20))
```

	634995	0	463	1	0.0	806.0	11.291044776119403	1.0	70.49513846124168	\
0	634995	0	463	1	0.0	806.0	11.291045	1.0	70.495138	
1	634995	0	463	1	0.0	806.0	11.291045	1.0	70.495138	
2	634995	0	463	1	0.0	806.0	11.291045	1.0	70.495138	
3	634995	0	463	1	0.0	806.0	11.291045	1.0	70.495138	
4	634995	0	463	1	0.0	806.0	11.291045	1.0	70.495138	
5	634995	0	463	1	0.0	806.0	11.291045	1.0	70.495138	

First problem :
columns are unnamed

Data overview

/04

Second major observation : there is no NaN or missing values :

```
▶ print (data.isnull().sum())
```

```
↳ 634995      0  
   0          0  
   463        0  
   1          0  
   0.0        0  
   806.0      0  
   11.291044776119403  0  
   1.0        0  
   70.49513846124168  0  
   0.0.1      0  
   806.0.1    0  
   7.574626865671642  0
```

but are all the 54 columns interesting?

Data overview

To solution this issue we did :

- spot the columns with irrelevant values with a dedicated algorithm
- plot the correlation matrix to find and drop unnecessary columns

-> 18 columns removed

issues encountered :

- the code used to drop every correlated column did not take account of the negative correlations
- one empty column was not detected by our algorithm
- some columns are correlated but still interesting (CC1 and CC4)

Selection of the output variable :




Figure 2. Demonstrating the essential feature details.

C1 was kept as the output feature

Creation of a model

/07

First observation : the dataset is divided in a test and a train part -> no need to make a pipeline

Results of the RandomForestRegressor without gridsearch :

```
Mean Squared Error: 780.8338210886157  
Best Hyperparameters: {'max_depth': 20, 'n_estimators': 200}  
R-squared: 0.9583827595059602
```

Now with gridsearch :

```
mse_best = mean_squared_error(y_test, y_pred_best)  
print(f"Mean Squared Error with Best Hyperparameters: {mse_best}")  
r2_best = r2_score(y_test, y_pred_best)  
print(f"R-squared: {r2_best}")
```

```
Mean Squared Error with Best Hyperparameters: 810.5752086252447  
R-squared: 0.9567975893400308
```

/08

Issue : the results obtained without gridsearch are better

Lets try Gradient Boosting Regressor

```
grid_search_gb = GridSearchCV(GradientBoostingRegressor(random_state=42), param_grid_gb, cv=5)
grid_search_gb.fit(X_train, y_train)
```

```
best_params_gb = grid_search_gb.best_params_
print(f"Best Hyperparameters (Gradient Boosting): {best_params_gb}")
r2_grad = r2_score(y_test, y_pred_gb)
print(f"R-squared: {r2_grad}")
```

```
Mean Squared Error (Gradient Boosting): 4878.267072022061
```

```
Best Hyperparameters (Gradient Boosting): {'learning_rate': 0.1, 'max_depth': 5, 'n_estimators': 50}
```

```
R-squared: 0.7399958756301656
```

We notice that we have pretty high MSE and a r2 closed to one. We may have only a few predictions failed, but that they are failed by a lot.

Creation of a model

/09

Bonus work predict CC2 in function of CC3 and CC4, let's try with random forest :

```
Mean Squared Error: 93.5792291889962  
Best Hyperparameters: {'max_depth': 10, 'n_estimators': 50}  
R-squared: 0.9842674800986632
```

Using gridsearch the MSE and R2 obtained are more than acceptable

Issue encountered : we tried to use a RandomForestClassifier to improve our model but the dataset is too big to be trained with this method

Conclusion

For our first model predicting cc1 in function of all the data available, we selected the **initial RandomForestRegressor**, giving us the following accuracy :

```
Mean Squared Error: 780.8338210886157
Best Hyperparameters: {'max_depth': 20, 'n_estimators': 200}
R-squared: 0.9583827595059602
```

For the bonus work predicting the number of comment obtained a day, just by looking at the number of comments the 2 days before we selected a **gridsearch optimized RandomForestRegressor** with the following accuracy :

```
Mean Squared Error: 93.5792291889962
Best Hyperparameters: {'max_depth': 10, 'n_estimators': 50}
R-squared: 0.9842674800986632
```

For both programs the results are very satisfying

Conclusion

